

CST 383 - Final Project

BayTech

- Warren Ngoun (wngoun@csumb.edu)
 - Yukio Rivera (yrivera@csumb.edu)
 - Jennah Yasin (jyasin@csumb.edu)
 - Luis Jimenez Barrios (ljimenezbarrios@csumb.edu)
-

Table of Contents:

- Intro
 - Our choice of dataset, What we're going to predict, features we plan to use as predictors
 - [Column Metadata](#)
 - [Imports](#)
 - [Data Preprocessing & Investigation](#)
 - [Data Encoding](#)
 - [Correlation and Heatmaps](#)
 - [Dealing with NaNs](#)
 - [Exploration & Visualization](#)
 - [Preliminary work on Machine Learning](#)
 - Choice of Classifier Model
 - Hyper-parameter Testing
 - Forward Feature Searching
 - Final Model results
 - [Conclusions](#)
-

Intro

[Back to Top](#)

We chose the Airline Passenger Satisfaction dataset for our final project because we are all interested in satisfying flight experiences and believe that it is important for airlines to take their passenger's reviews into consideration.

- Link to Kaggle Repo: <https://www.kaggle.com/datasets/teejmahal20/airline-passenger-satisfaction>

With this dataset, we are going to predict what factors may be most relevant and most correlated to the passenger's satisfaction. We will be analyzing which features are most correlated to our topic and dropping those that aren't as relevant/needed.

Also a note, the file was pre-processed by our Kaggle author, he split the original data set into two separate files (test and train) while also removing a vast majority of NaN and bad data, so most of that work was already done for us by the repo author.

We plan to use the "satisfaction" feature as our target label for this project. The initial features that we plan to use as predictors for our project are: Seat comfort, in-flight entertainment, cleanliness, and food & drinks, but that may change as we test them out later.

Column Contents

Taken from the Kaggle Repo description.

- Gender: Gender of the passengers (Female, Male)
- Customer Type: The customer type (Loyal customer, disloyal customer)
- Age: The actual age of the passengers
- Type of Travel: Purpose of the flight of the passengers (Personal Travel, Business Travel)
- Class: Travel class in the plane of the passengers (Business, Eco, Eco Plus)
- Flight distance: The flight distance of this journey
- Inflight wifi service: Satisfaction level of the inflight wifi service (0: Not Applicable; 1-5 stars)
- Departure/Arrival time convenient: Satisfaction level of Departure/Arrival time convenient
- Ease of Online booking: Satisfaction level of online booking
- Gate location: Satisfaction level of Gate location
- Food and drink: Satisfaction level of Food and drink
- Online boarding: Satisfaction level of online boarding (0: Not Applicable; 1-5 stars)
- Seat comfort: Satisfaction level of Seat comfort
- Inflight entertainment: Satisfaction level of inflight entertainment
- On-board service: Satisfaction level of On-board service
- Leg room service: Satisfaction level of Leg room service
- Baggage handling: Satisfaction level of baggage handling
- Check-in service: Satisfaction level of Check-in service
- Inflight service: Satisfaction level of inflight service
- Cleanliness: Satisfaction level of Cleanliness
- Departure Delay in Minutes: Minutes delayed when departure
- Arrival Delay in Minutes: Minutes delayed when Arrival
- Satisfaction: Airline satisfaction level(Satisfaction or 0 , neutral/dissatisfied or 1)

Imports

All the necessary imports we need for the project.

```
In [625.. import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
# For ML Work
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import learning_curve
from sklearn.metrics import confusion_matrix
```

From our dataset, we have two files from the original Kaggle repo, an already split training and test csv,

In [626...

In [627...

In [628]:

Out[628]:

[illegible]

mean	51951.5	64924.2	39.4	1189.4	2.7	3.1	2.8	3.0	3.
std	29994.6	37463.8	15.1	997.1	1.3	1.5	1.4	1.3	1.
min	0.0	1.0	7.0	31.0	0.0	0.0	0.0	0.0	0.
25%	25975.8	32533.8	27.0	414.0	2.0	2.0	2.0	2.0	2.
50%	51951.5	64856.5	40.0	843.0	3.0	3.0	3.0	3.0	3.
75%	77927.2	97368.2	51.0	1743.0	4.0	4.0	4.0	4.0	4.
max	103903.0	129880.0	85.0	4983.0	5.0	5.0	5.0	5.0	5.

Lets see which columns are categorical.

```
In [629... df.dtypes[df.dtypes == 'object']
```

```
Out[629]: Gender          object
Customer Type    object
Type of Travel   object
Class            object
satisfaction     object
dtype: object
```

For each of those above categories, lets investigate the values.

```
In [630... print(df['Gender'].value_counts())
```

```
Female    52727
Male      51177
Name: Gender, dtype: int64
```

```
In [631... print(df['Customer Type'].value_counts())
```

```
Loyal Customer    84923
disloyal Customer  18981
Name: Customer Type, dtype: int64
```

```
In [632... print(df['Type of Travel'].value_counts())
```

```
Business travel    71655
Personal Travel    32249
Name: Type of Travel, dtype: int64
```

```
In [633... print(df['Class'].value_counts())
```

```
Business    49665
Eco         46745
Eco Plus    7494
Name: Class, dtype: int64
```

```
In [634... print(df['Type of Travel'].value_counts())
```

```
Business travel    71655
Personal Travel    32249
Name: Type of Travel, dtype: int64
```

Data Encoding

Now we'll go ahead and One-Hot Encode all of those categorical values and then label encode *satisfaction* so we can use a single column for the [machine learning](#).

```
In [635... cols = ['Gender', 'Type of Travel', 'Class', 'Customer Type']
# Keeping these as backups for graphing later.
```

```

classForGraphing = df['Class']
genderForGraphing = df['Gender']

for col in cols:
    catCol = pd.get_dummies(df[col], prefix=col)
    df.drop(col, axis=1, inplace=True)
    df = pd.concat([df, catCol], axis=1)
    # Repeat but for our test data set too.
    catCol2 = pd.get_dummies(dfTest[col], prefix=col)
    dfTest.drop(col, axis=1, inplace=True)
    dfTest = pd.concat([dfTest, catCol2], axis=1)

```

```

In [636.. # Here we made a copy so that when we later try and visualize the data we can get proper
satisfactionForGraphing = df['satisfaction'].copy()
# Checking before and after. 0 = neutral/dissatisfied 1 = satisfied
print(df['satisfaction'].value_counts())
df['satisfaction'] = df['satisfaction'].astype('category').cat.codes
print(df['satisfaction'].value_counts())

```

```

neutral or dissatisfied    58879
satisfied                  45025
Name: satisfaction, dtype: int64
0      58879
1      45025
Name: satisfaction, dtype: int64

```

Lets check out the different columns from both of our data frames now.

```

In [637.. df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103904 entries, 0 to 103903
Data columns (total 30 columns):
 #   Column                                                                 Non-Null Count  Dtype
---  -
 0   Unnamed: 0                                                            103904 non-null  int64
 1   id                                                                    103904 non-null  int64
 2   Age                                                                    103904 non-null  int64
 3   Flight Distance                                                       103904 non-null  int64
 4   Inflight wifi service                                                 103904 non-null  int64
 5   Departure/Arrival time convenient 103904 non-null  int64
 6   Ease of Online booking                                                103904 non-null  int64
 7   Gate location                                                         103904 non-null  int64
 8   Food and drink                                                        103904 non-null  int64
 9   Online boarding                                                       103904 non-null  int64
10   Seat comfort                                                          103904 non-null  int64
11   Inflight entertainment                                                103904 non-null  int64
12   On-board service                                                      103904 non-null  int64
13   Leg room service                                                      103904 non-null  int64
14   Baggage handling                                                      103904 non-null  int64
15   Checkin service                                                       103904 non-null  int64
16   Inflight service                                                      103904 non-null  int64
17   Cleanliness                                                           103904 non-null  int64
18   Departure Delay in Minutes                                            103904 non-null  int64
19   Arrival Delay in Minutes                                             103594 non-null  float64
20   satisfaction                                                           103904 non-null  int8
21   Gender_Female                                                         103904 non-null  uint8
22   Gender_Male                                                           103904 non-null  uint8
23   Type of Travel_Business travel 103904 non-null  uint8
24   Type of Travel_Personal Travel 103904 non-null  uint8
25   Class_Business                                                        103904 non-null  uint8
26   Class_Eco                                                            103904 non-null  uint8
27   Class_Eco Plus                                                       103904 non-null  uint8
28   Customer Type_Loyal Customer 103904 non-null  uint8
29   Customer Type_disloyal Customer 103904 non-null  uint8

```

```
dtypes: float64(1), int64(19), int8(1), uint8(9)
memory usage: 16.8 MB
```

```
In [638.. dfTest.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25976 entries, 0 to 25975
Data columns (total 30 columns):
#   Column                                          Non-Null Count  Dtype
---  -
0   Unnamed: 0                                     25976 non-null  int64
1   id                                              25976 non-null  int64
2   Age                                             25976 non-null  int64
3   Flight Distance                               25976 non-null  int64
4   Inflight wifi service                         25976 non-null  int64
5   Departure/Arrival time convenient             25976 non-null  int64
6   Ease of Online booking                       25976 non-null  int64
7   Gate location                                 25976 non-null  int64
8   Food and drink                               25976 non-null  int64
9   Online boarding                              25976 non-null  int64
10  Seat comfort                                  25976 non-null  int64
11  Inflight entertainment                       25976 non-null  int64
12  On-board service                             25976 non-null  int64
13  Leg room service                             25976 non-null  int64
14  Baggage handling                             25976 non-null  int64
15  Checkin service                             25976 non-null  int64
16  Inflight service                             25976 non-null  int64
17  Cleanliness                                  25976 non-null  int64
18  Departure Delay in Minutes                   25976 non-null  int64
19  Arrival Delay in Minutes                     25893 non-null  float64
20  satisfaction                                  25976 non-null  object
21  Gender_Female                                25976 non-null  uint8
22  Gender_Male                                  25976 non-null  uint8
23  Type of Travel_Business travel               25976 non-null  uint8
24  Type of Travel_Personal Travel               25976 non-null  uint8
25  Class_Business                               25976 non-null  uint8
26  Class_Eco                                    25976 non-null  uint8
27  Class_Eco Plus                               25976 non-null  uint8
28  Customer Type_Loyal Customer                 25976 non-null  uint8
29  Customer Type_disloyal Customer              25976 non-null  uint8
dtypes: float64(1), int64(19), object(1), uint8(9)
memory usage: 4.4+ MB
```

We went from 24 columns to 29 now after encoding.

NaN Processing

Now we're interested if there are still any NaN values left in our dataset and then if the data will be suitable for imputation.

```
In [639.. print(f'There are {df.isna().sum().sum()} NaN values in df.')
df.isna().sum()
```

```
There are 310 NaN values in df.
```

```
Out[639]: Unnamed: 0      0
id            0
Age           0
Flight Distance  0
Inflight wifi service  0
Departure/Arrival time convenient  0
Ease of Online booking  0
Gate location  0
Food and drink  0
Online boarding  0
```

Seat comfort	0
Inflight entertainment	0
On-board service	0
Leg room service	0
Baggage handling	0
Checkin service	0
Inflight service	0
Cleanliness	0
Departure Delay in Minutes	0
Arrival Delay in Minutes	310
satisfaction	0
Gender_Female	0
Gender_Male	0
Type of Travel_Business travel	0
Type of Travel_Personal Travel	0
Class_Business	0
Class_Eco	0
Class_Eco Plus	0
Customer Type_Loyal Customer	0
Customer Type_disloyal Customer	0
dtype: int64	

```
In [640]: print(f'There are {dfTest.isna().sum().sum()} NaN values in our test df.')
dfTest.isna().sum()
```

There are 83 NaN values in our test df.

```
Out[640]: Unnamed: 0      0
id      0
Age      0
Flight Distance      0
Inflight wifi service      0
Departure/Arrival time convenient      0
Ease of Online booking      0
Gate location      0
Food and drink      0
Online boarding      0
Seat comfort      0
Inflight entertainment      0
On-board service      0
Leg room service      0
Baggage handling      0
Checkin service      0
Inflight service      0
Cleanliness      0
Departure Delay in Minutes      0
Arrival Delay in Minutes      83
satisfaction      0
Gender_Female      0
Gender_Male      0
Type of Travel_Business travel      0
Type of Travel_Personal Travel      0
Class_Business      0
Class_Eco      0
Class_Eco Plus      0
Customer Type_Loyal Customer      0
Customer Type_disloyal Customer      0
dtype: int64
```

```
In [641]: # Calculate the average of the column
average_delay = df['Arrival Delay in Minutes'].mean()
average_delay_test = dfTest['Arrival Delay in Minutes'].mean()
# Impute & replace NaNs with the average value
df['Arrival Delay in Minutes'].fillna(value=average_delay, inplace=True)
dfTest['Arrival Delay in Minutes'].fillna(value=average_delay_test, inplace=True)
```

```
In [642... print(f'There are {df.isna().sum().sum()} NaN values in df.')
print(f'There are {dfTest.isna().sum().sum()} NaN values in our test df.')
```


There are 0 NaN values in df.
There are 0 NaN values in our test df.

Correlation Heatmap & Column Dropping

We're interested in also trimming out unnecessary columns from the dataset so there are less potential noisy predictors. Here we have two correlation heat maps, one of before and after a purging of potential columns.


```
In [643... corrDfOld = df.corr()
sns.heatmap(corrDfOld, xticklabels=corrDfOld.columns, yticklabels=corrDfOld.columns)
```

Out[643]: <AxesSubplot:>



```
In [644... corrDf = df.copy()
corrDf.drop(['id', 'Unnamed: 0'], axis=1, inplace=True)
correlation = corrDf.corr()
sns.heatmap(correlation, xticklabels=correlation.columns, yticklabels=correlation.columns)
```

Out[644]: <AxesSubplot:>



There weren't too many hard correlations in our data now after encoding, other than some with the data formerly in encoded in binary categories like *Gender*. So considering that, we officially drop the "extra" columns from our dataset and investigate what's left again, leaving a majority of the data intact after all. We are now ultimately left with 27 columns, 26 potential predictors and 1 label.

```
In [645.. toDropCols = ['id', 'Unnamed: 0']
df.drop(toDropCols, axis=1, inplace=True)
dfTest.drop(toDropCols, axis=1, inplace=True)
```

```
In [646.. df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 103904 entries, 0 to 103903
Data columns (total 28 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age                                  103904 non-null  int64
1   Flight Distance                     103904 non-null  int64
2   Inflight wifi service               103904 non-null  int64
3   Departure/Arrival time convenient  103904 non-null  int64
4   Ease of Online booking              103904 non-null  int64
5   Gate location                      103904 non-null  int64
6   Food and drink                     103904 non-null  int64
7   Online boarding                    103904 non-null  int64
8   Seat comfort                       103904 non-null  int64
9   Inflight entertainment              103904 non-null  int64
```

```
10 On-board service 103904 non-null int64
11 Leg room service 103904 non-null int64
12 Baggage handling 103904 non-null int64
13 Checkin service 103904 non-null int64
14 Inflight service 103904 non-null int64
15 Cleanliness 103904 non-null int64
16 Departure Delay in Minutes 103904 non-null int64
17 Arrival Delay in Minutes 103904 non-null float64
18 satisfaction 103904 non-null int8
19 Gender_Female 103904 non-null uint8
20 Gender_Male 103904 non-null uint8
21 Type of Travel_Business travel 103904 non-null uint8
22 Type of Travel_Personal Travel 103904 non-null uint8
23 Class_Business 103904 non-null uint8
24 Class_Eco 103904 non-null uint8
25 Class_Eco Plus 103904 non-null uint8
26 Customer Type_Loyal Customer 103904 non-null uint8
27 Customer Type_disloyal Customer 103904 non-null uint8
dtypes: float64(1), int64(17), int8(1), uint8(9)
memory usage: 15.3 MB
```

Data Exploration

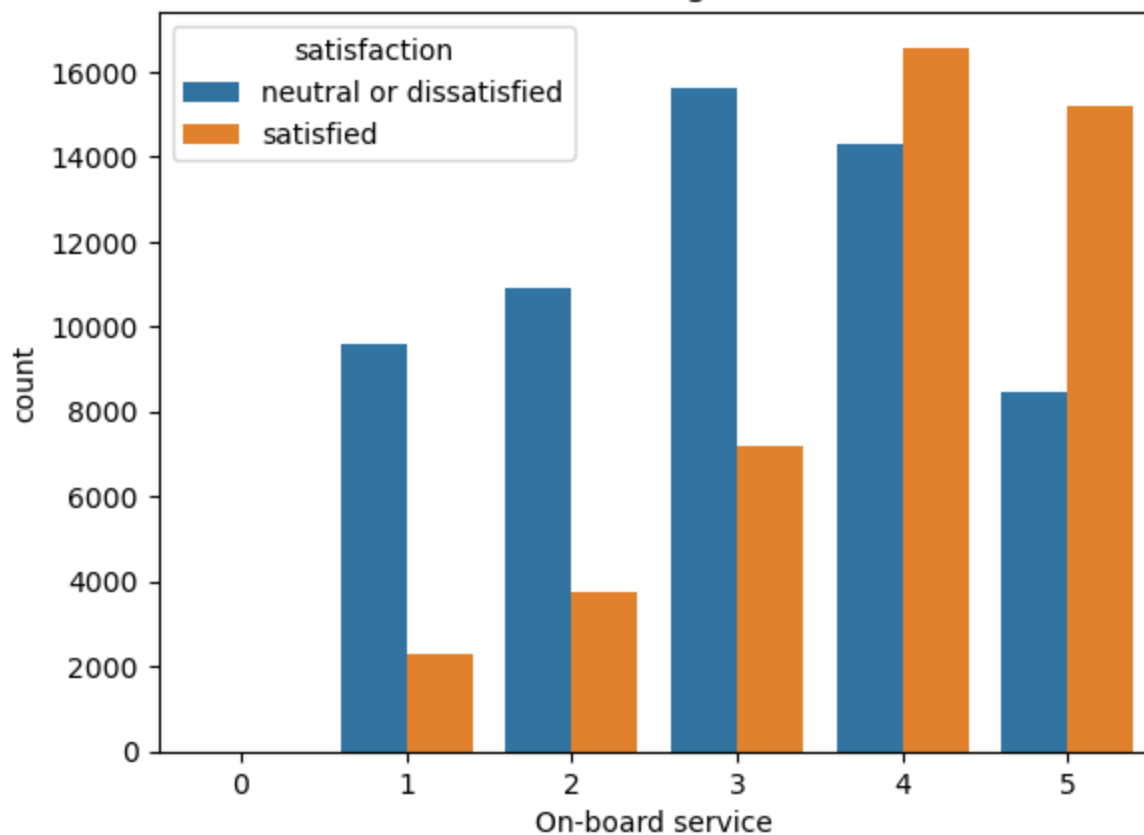
[Back to Top](#)

Now we look at certain columns and the relationships between them via graphing. Our satisfaction (which will be our labels) are a binary option with only two answers (neutral or dissatisfied or satisfied).

Visualizations


```
In [647... # Passenger's satisfaction based on the Ratings of on-board service.
sns.countplot(x=df['On-board service'], hue=satisfactionForGraphing)
plt.title('Satisfaction based on Ratings for On-board Service')
plt.show()
```

Satisfaction based on Ratings for On-board Service



In [648...

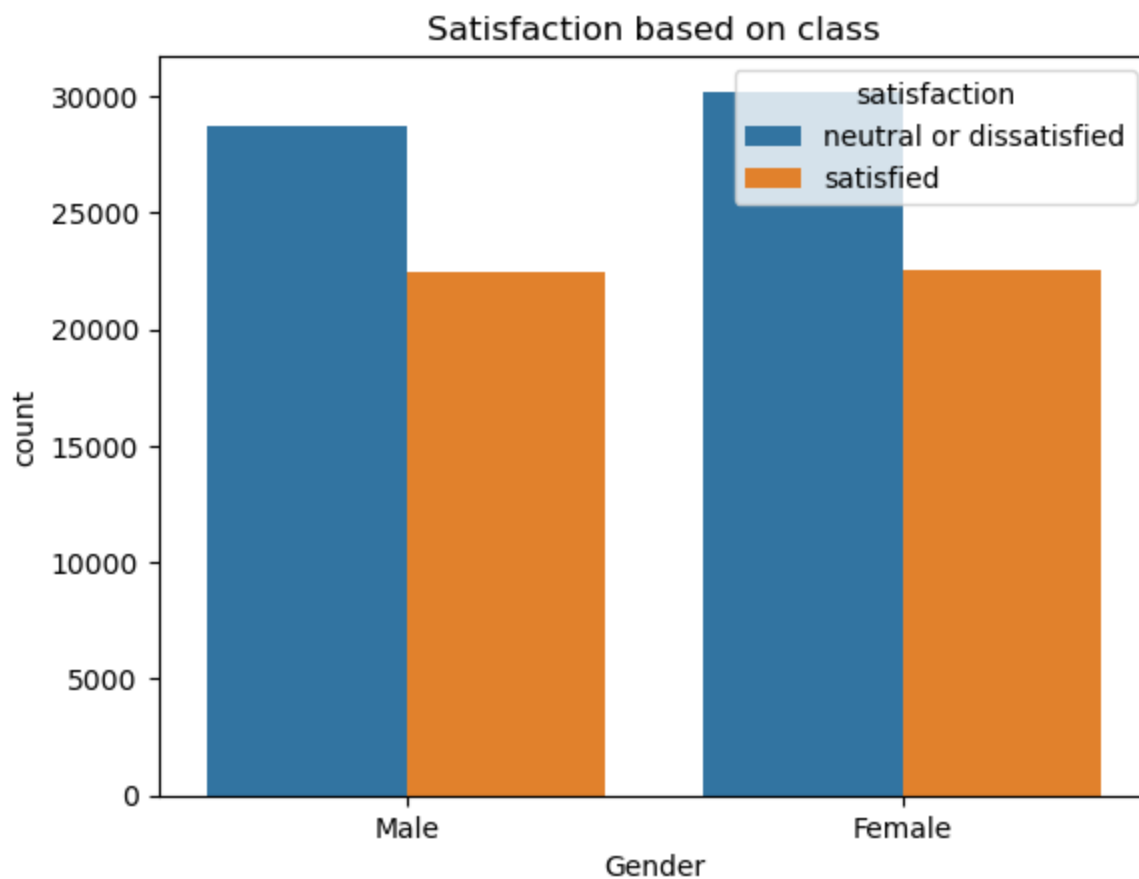
```
# Passenger's satisfaction based off of Class
sns.countplot(x=classForGraphing, hue=satisfactionForGraphing)
plt.title('Satisfaction based on class')
plt.show()
```



In [649...


```
# Passenger's satisfaction based off of their gender. To see if the graphs are near equa
sns.countplot(x=genderForGraphing, hue=satisfactionForGraphing)
```

```
plt.title('Satisfaction based on class')  
plt.show()
```



```
In [650... # Passenger's Satisfaction based off of Seat Comfort Ratings  
sns.boxplot(x=df['Seat comfort'], y=satisfactionForGraphing)  
plt.title('Satisfaction based on Seat Comfort')  
plt.xlabel('Seat Comfort Rating (1-5)')  
plt.ylabel('Satisfaction')
```

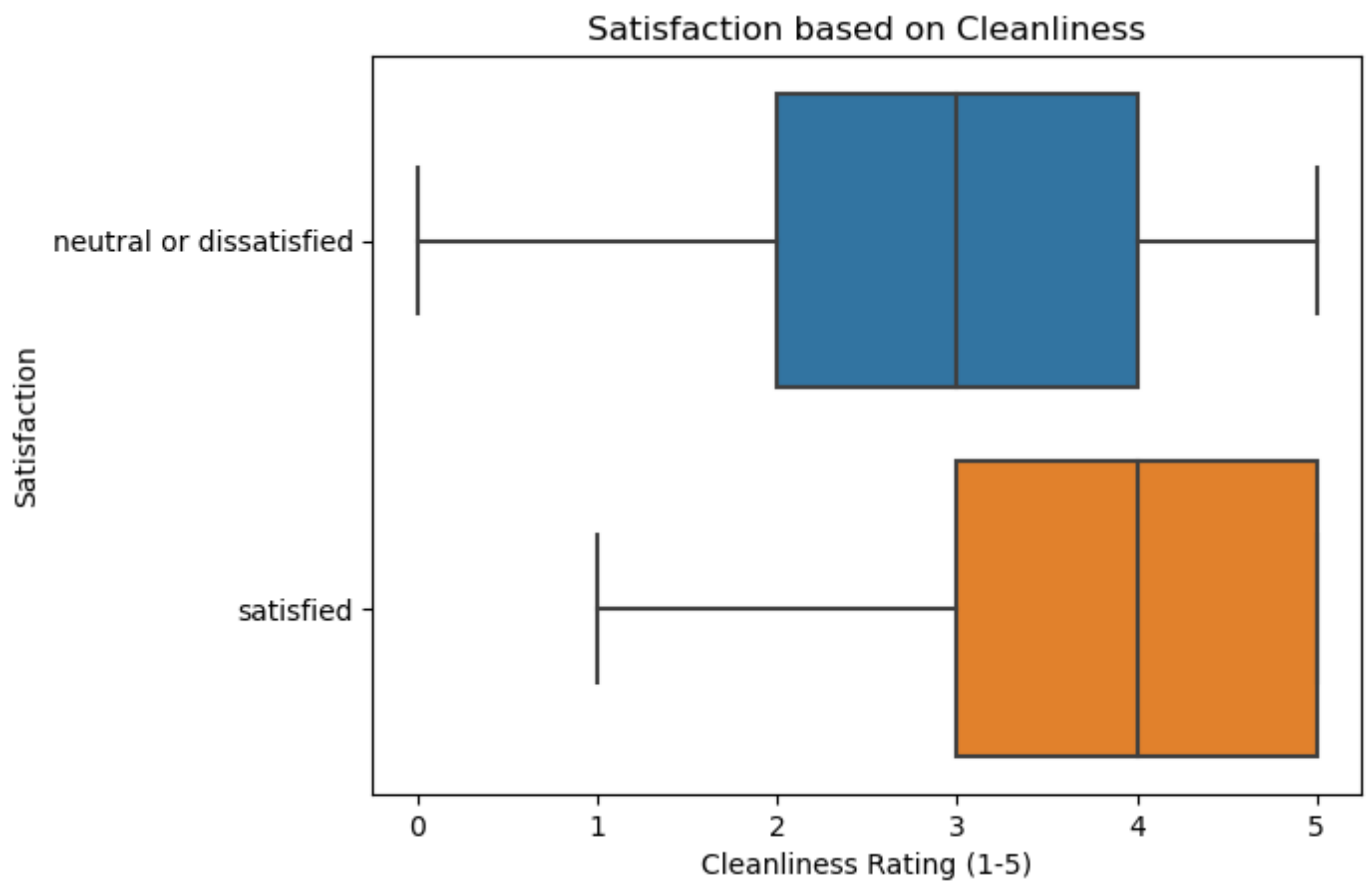
```
Out[650]: Text(0, 0.5, 'Satisfaction')
```



Based on the results, it appears that passengers who rated Seat Comfort between a 4-5 (high comfort) were satisfied compared to those who rated the seat comfort less than a 4, were not satisfied. There are two outliers that have chosen to be satisfied although they rated the seat comfort a 1 or 2.

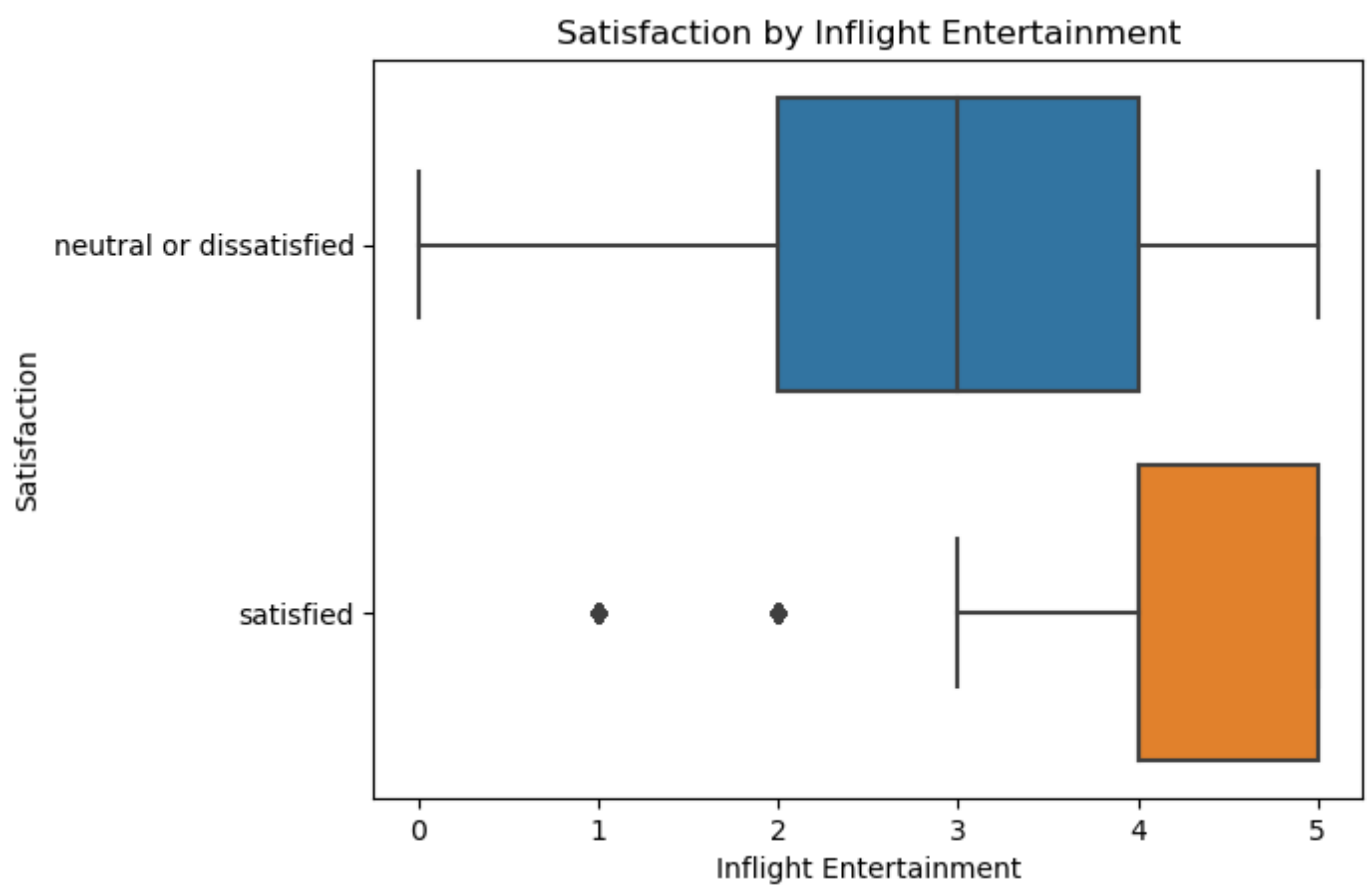
```
In [651]: # Passenger's Satisfaction based off of Cleanliness Ratings
sns.boxplot(data=df, x=df['Cleanliness'], y=satisfactionForGraphing)
plt.xlabel('Cleanliness Rating (1-5)')
plt.ylabel('Satisfaction')
plt.title('Satisfaction based on Cleanliness')
```

```
Out[651]: Text(0.5, 1.0, 'Satisfaction based on Cleanliness')
```

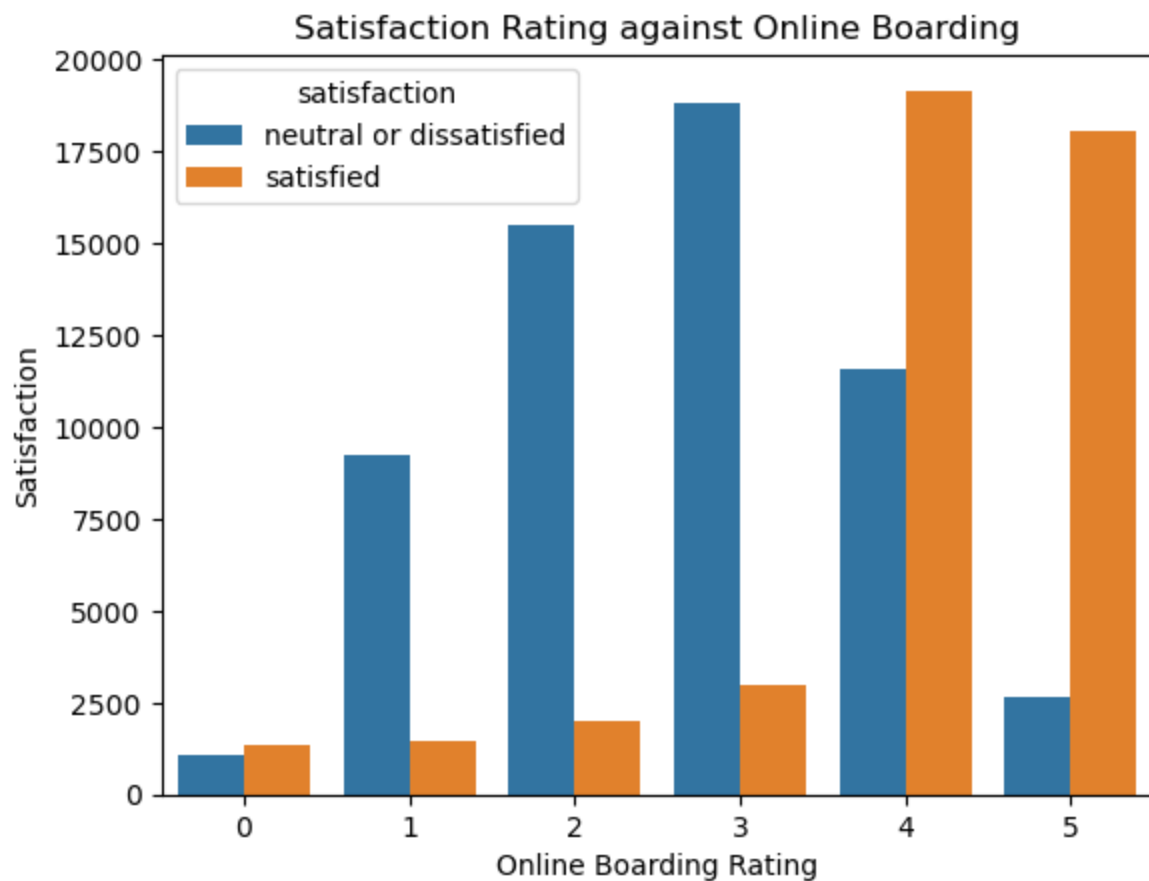


This plot shows the higher the cleanliness rating (3-5), the more likely satisfied the passenger was and the lower the cleanliness rating(1~3), the more likely they weren't satisfied.

```
In [652... # Passenger's Satisfaction based off of Inflight Entertainment
sns.boxplot(x='Inflight entertainment', y=satisfactionForGraphing, data=df)
plt.title('Satisfaction by Inflight Entertainment')
plt.xlabel('Inflight Entertainment')
plt.ylabel('Satisfaction')
plt.show()
```

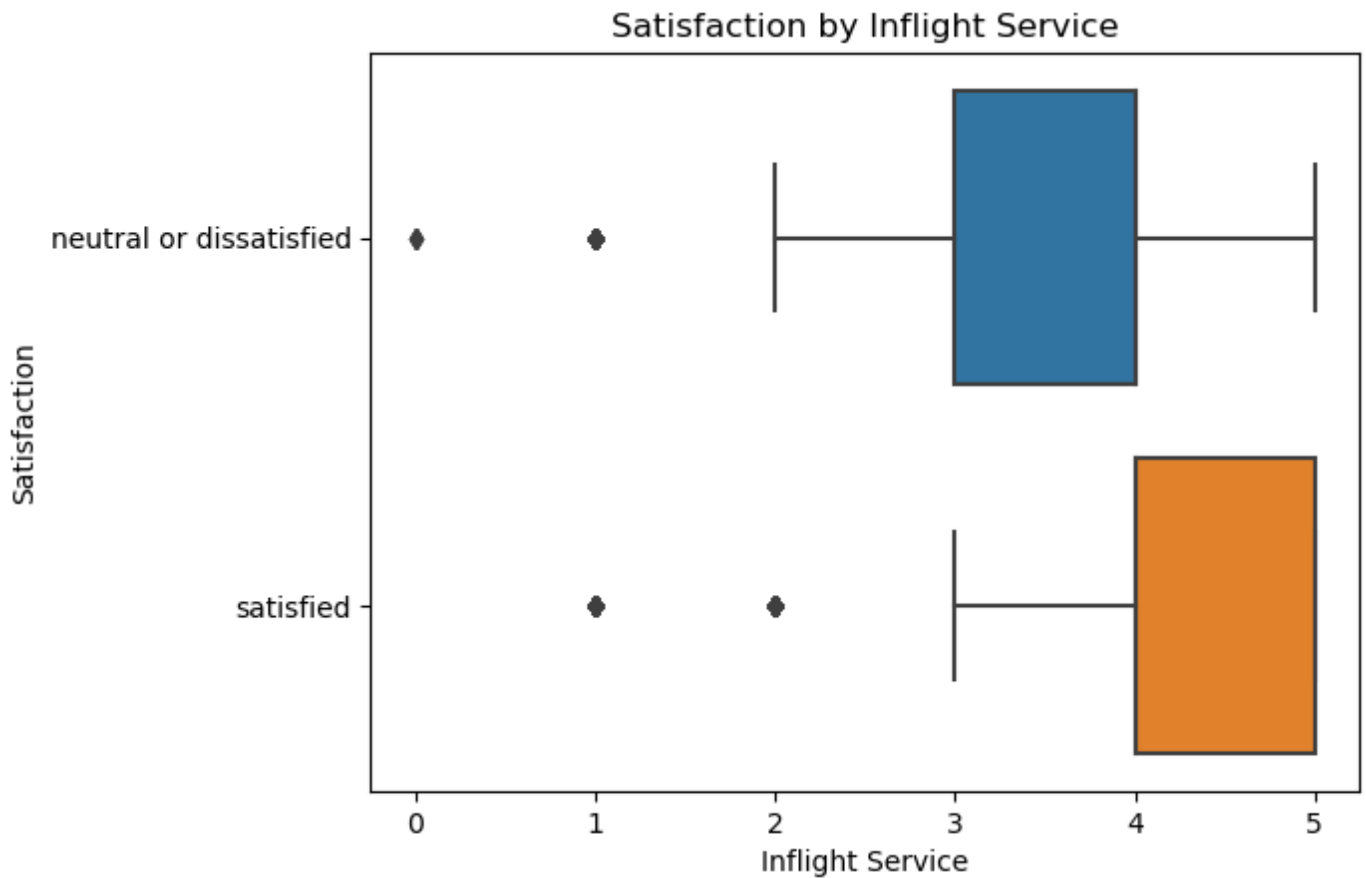


```
In [653... # Checking overall satisfaction for all online boarders.
sns.countplot(x='Online boarding', hue=satisfactionForGraphing, data=df)
plt.title('Satisfaction Rating against Online Boarding')
plt.xlabel('Online Boarding Rating')
plt.ylabel('Satisfaction')
plt.show()
```



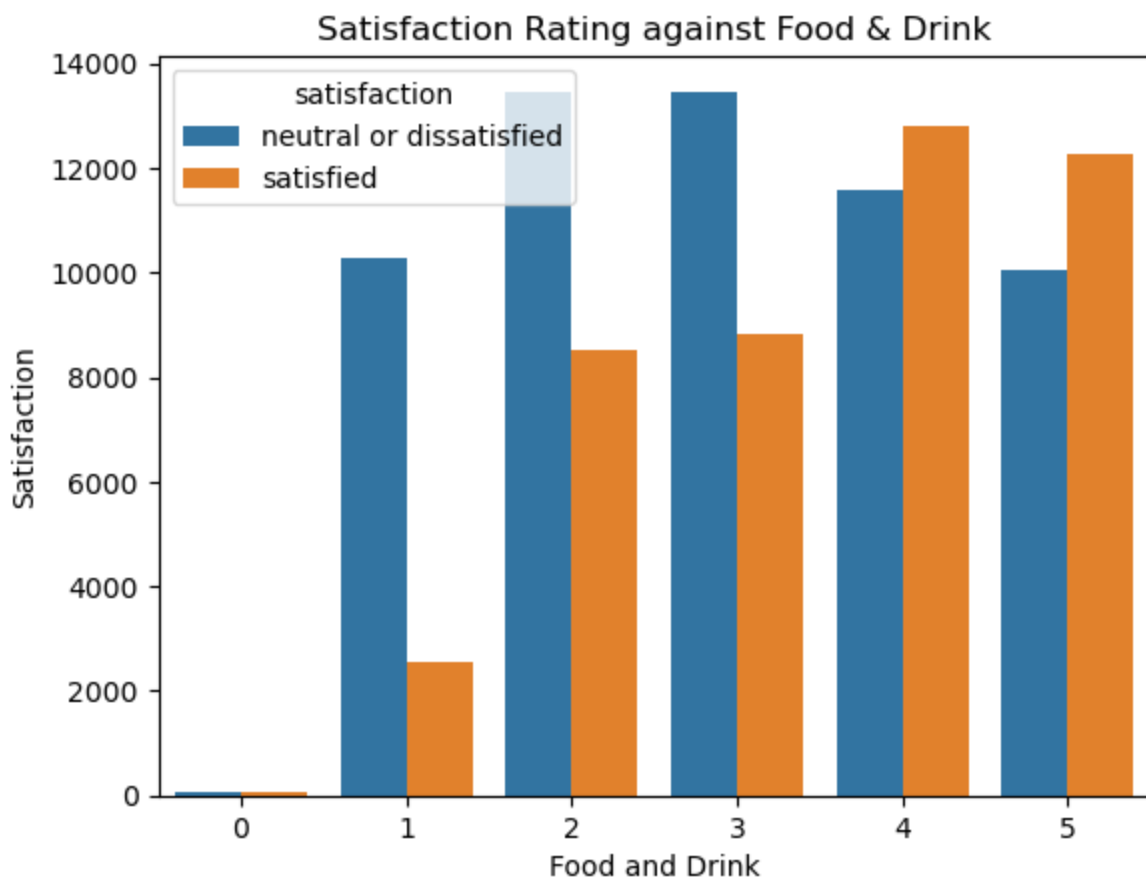
In [654...

```
# Passenger's Satisfaction based off of Inflight Service
sns.boxplot(x='Inflight service', y=satisfactionForGraphing, data=df)
plt.title('Satisfaction by Inflight Service')
plt.xlabel('Inflight Service')
plt.ylabel('Satisfaction')
plt.show()
```



In [655...

```
sns.countplot(x='Food and drink', hue=satisfactionForGraphing, data=df)
plt.title('Satisfaction Rating against Food & Drink')
plt.xlabel('Food and Drink')
plt.ylabel('Satisfaction')
plt.show()
```

Training

[Back to Top](#)

Since we have only two results in our *satisfaction* column (neutral or dissatisfied & satisfied), we are going to use either kNN Classifier or a Decision Tree Classifier for our project.

To prep for the actual machine learning, we'll prep our *y_train* & *y_test* so that we can reuse them in all subsequent trials without having to redeclare/tweak it.

```
In [656... y_train = df['satisfaction'].values
y_test = dfTest['satisfaction'].astype('category').cat.codes
```

A helper function that simulates *train_test_split* but also with the added feature of normalizing our data. Also only works for X since we already did Y earlier.

```
In [657... def x_train_test_split(predictors):
    scaler = StandardScaler()
    xPred = scaler.fit_transform(df[predictors].values)
    testPred = scaler.fit_transform(dfTest[predictors].values)
    return (xPred, testPred)
```

Also a note for these future cells, some of them take a generous amount of computational time due to large amounts of data needing to be processed.

```
In [658... predictors = ['Inflight entertainment', 'On-board service', 'Leg room service', 'Baggage
X_train, X_test = x_train_test_split(predictors)
```

```
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)

predictions = knn.predict(X_test)
accuracy = (predictions == y_test).mean()
print(f'kNN 5-feature accuracy: {accuracy.round(2)}')
```

```
c:\Users\warre\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
  mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
kNN 5-feature accuracy: 0.73
```

Now that we have a simple kNN Classifier model setup, we decided to run another test with all the columns as potential predictors. We also decided to tweak the hyper-parameters to see if we could speed up the kNN model.

```
In [659... dfCopy = df.copy()
dfCopy.drop('satisfaction', inplace=True, axis=1)
predictors = dfCopy.columns
X_train, X_test = x_train_test_split(predictors)

knn = KNeighborsClassifier(algorithm='brute') # 25s
# knn = KNeighborsClassifier(algorithm='ball_tree') # over a min
# knn = KNeighborsClassifier(algorithm='kd_tree') # 40s
knn.fit(X_train, y_train)

predictions = knn.predict(X_test)
accuracy = (predictions == y_test).mean()
print(f'kNN all-feature accuracy: {accuracy.round(2)}')
```

```
c:\Users\warre\anaconda3\lib\site-packages\sklearn\neighbors\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.
  mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
kNN all-feature accuracy: 0.93
```

```
In [660... dfCopy = df.copy()
dfCopy.drop('satisfaction', inplace=True, axis=1)
predictors = dfCopy.columns
X_train, X_test = x_train_test_split(predictors)

# tree = DecisionTreeClassifier(max_depth=2) # 0.86
tree = DecisionTreeClassifier(max_depth=4) # 0.89, close enough to kNN accuracy but keep
tree.fit(X_train, y_train)

predictions = tree.predict(X_test)
accuracy = (predictions == y_test).mean()
print(f'Decision Tree all-feature accuracy: {accuracy.round(2)}')
```

```
Decision Tree all-feature accuracy: 0.89
```

Considering the lengthy runtime of kNN compared to the decision trees (30-40s vs ~0.3s) we decided to use **Decision Trees** from now on as our main model going forward.

We'll start by investigating hyper parameters to improve our models accuracy. First, we'll tweak the **max_depth** and create a learning curve graph to pick out the optimal one.

```

In [661... # Code snippet from previous lab. Code to run plenty of tests with our model that'll map
k = 10
overallTE = []
overallTR = []
dfCopy = df.copy()
dfCopy.drop('satisfaction', inplace=True, axis=1)
predictors = dfCopy.columns
X_train, X_test = x_train_test_split(predictors)

for i in range(1, k+1, 1):
    knn = DecisionTreeClassifier(max_depth=i)
    te_errs = []
    tr_errs = []
    tr_sizes = np.linspace(100, X_train.shape[0], 10).astype(int)
    for tr_size in tr_sizes:
        # train model on a subset of the training data
        X_train1 = X_train[:tr_size,:]
        y_train1 = y_train[:tr_size]
        knn.fit(X_train1, y_train1)
        # Errors from Training & Test Data
        tr_predicted = knn.predict(X_train1)
        err = (tr_predicted != y_train1).mean()
        tr_errs.append(err)
        te_predicted = knn.predict(X_test)
        err = (te_predicted != y_test).mean()
        te_errs.append(err)
        # Calc the learning curve values and append them for later.
    tr_sizes, tr_errs, te_errs = learning_curve(
        knn, X_train, y_train, cv=10, scoring='accuracy')
    overallTR.append(np.mean(tr_errs, axis=1))
    overallTE.append(np.mean(te_errs, axis=1))


```

```

In [662... # Same snippet from lab, but separated for easier tweaking of graphs.
# Make the resulting pairs "easier" to interpret
color = ['red', 'green', 'blue', 'orange', 'purple', 'gold', 'violet', 'maroon', 'pink',
k = 1
for i in range(0, len(overallTE), 1):
    plt.plot(tr_sizes, overallTE[i], label=f'test {k}', color=color[i], ls='dashdot')
    plt.plot(tr_sizes, overallTR[i], label=f'train {k}', color=color[i], linewidth=1.5,
    k = k+1
plt.legend(loc='right', prop={'size': 8})
plt.xlabel('Training Sizes')
plt.ylabel('Accuracy')
plt.title('Learning Curve of Decision Tree w/ Different Max Depths')
plt.show()

```

Learning Curve of Decision Tree w/ Different Max Depths



From this graph, we can see all the different gaps between the `max_depth` values and how the overall accuracy goes up as the depth increases. This can definitely be over fit if we go too far up, so we'll pick something in the middle with a small gap and then use that parameter for subsequent tests.

So going forward, `max_depth` will be set to 4.

We had also experimented with `GridSearchCV` to test out potential hyper-parameter improvements, but not only did the cell take around a minute to run, the results didn't improve our prediction accuracy much so we stuck with the depth parameter only.

```
In [663... dfCopy = df.copy()
dfCopy.drop('satisfaction', inplace=True, axis=1)
predictors = dfCopy.columns
X_train, X_test = x_train_test_split(predictors)

parameters = [{'min_samples_leaf': [0.1, 0.2, 0.3], 'max_leaf_nodes': [4, 8, 16]}]
tree = DecisionTreeClassifier(max_depth=4)
test = GridSearchCV(tree, parameters, scoring='accuracy', cv=10)
test.fit(X_train, y_train)

print(f'Our best score was: {test.best_score_} and the best params were {test.best_param

Our best score was: 0.8435478840845556 and the best params were {'max_leaf_nodes': 4, 'm
in_samples_leaf': 0.1}.
```

Now we have figured out our major hyper-parameter and can begin checking for the best predictors/features.

Finding best feature that has highest accuracy:

```
In [664... dfCopy = df.copy()
dfCopy.drop(columns='satisfaction', axis=1, inplace=True)
predictors = dfCopy.columns
```

```

# Was unable to setup the proper file split here for some reason.
X_train_feat, X_test_feat, y_train_feat, y_test_feat = train_test_split(dfCopy, df['sati

colName = []
currentAccuracy = 0
for col in predictors:
    X_train_1 = X_train_feat[[col]]
    scores = cross_val_score(DecisionTreeClassifier(random_state = 42), X_train_1, y_tra
    accuracy = scores.mean()
    if (accuracy > currentAccuracy):
        currentAccuracy = accuracy
        colName = col
print('Best Feature: {}, Best Accuracy: {:.2f}%'.format(colName, currentAccuracy))

```

Best Feature: Online boarding, Best Accuracy: 0.79%

Top 10 combined features that have highest accuracy using forward feature search:

```

In [665... dfCopy = df.copy()
dfCopy.drop(columns='satisfaction', axis=1, inplace=True)
predictors = dfCopy.columns
# Same as unable to setup the proper file split, but its fine for training
X_train_feat, X_test_feat, y_train_feat, y_test_feat = train_test_split(dfCopy, df['sati

remaining = list(predictors)
selected = []
n = 10

while len(selected) < n:
    currentAccuracy = 0
    colName = ''
    for feature in remaining:
        X_selected = X_train_feat[selected + [feature]]
        scores = cross_val_score(DecisionTreeClassifier(random_state = 42), X_selected,
        accuracy = scores.mean()
        if (accuracy > currentAccuracy):
            currentAccuracy = accuracy
            colName = feature

    remaining.remove(colName)
    selected.append(colName)
    print('Feature: {}, Accuracy: {:.2f}'.format(colName, currentAccuracy))

```

```

Feature: Online boarding, Accuracy: 0.79
Feature: Type of Travel_Business travel, Accuracy: 0.85
Feature: Inflight wifi service, Accuracy: 0.89
Feature: Gate location, Accuracy: 0.92
Feature: Baggage handling, Accuracy: 0.93
Feature: Customer Type_disloyal Customer, Accuracy: 0.94
Feature: Class_Business, Accuracy: 0.95
Feature: Inflight service, Accuracy: 0.95
Feature: Seat comfort, Accuracy: 0.95
Feature: Customer Type_Loyal Customer, Accuracy: 0.95

```

From the previous cell, we determined we only need around roughly ~90% accuracy, so we decided to use just the top 5 features instead of all 10.

```

In [666... predictors = ['Online boarding', 'Type of Travel_Business travel', 'Inflight wifi servic
# True file split was achieved for the final model run
X_train, X_test = x_train_test_split(predictors)

tree = DecisionTreeClassifier(max_depth=4, random_state=42)
tree.fit(X_train, y_train)

predictions = tree.predict(X_test)

```

```
accuracy = (predictions == y_test).mean()
print(f'Final Decision Tree Accuracy: {accuracy.round(2)}%')
```

Final Decision Tree Accuracy: 0.88%

Conclusions

[Back to Top](#)

Here we test our model by adding in parameters (the predictors) that signaled a **dissatisfied** customer (Arrival Delay in Minutes, Baggage Handling, etc.) to see if our model would most predict an unhappy customer. We did the same with **satisfied** and got the output we expected.

```
In [667.. # Create a dictionary with the feature values for a single demo customer.
demoCustomer = {'Gender_Female': 0, 'Gender_Male': 1,
                'Type of Travel_Business travel': 0,
                'Type of Travel_Personal Travel': 0, 'Class_Business': 0, 'Class_Eco': 1,
                'Class_Eco Plus': 0, 'Customer Type_Loyal Customer': 1, 'Customer Type_dislo
                'Age': 35, 'Flight Distance': 1000, 'Inflight wifi service': 1, 'Departure/A
                'Ease of Online booking': 4, 'Gate location': 1, 'Food and drink': 4, 'Onlin
                'Seat comfort': 3, 'Inflight entertainment': 4, 'On-board service': 4, 'Leg
                'Baggage handling': 0, 'Checkin service': 5, 'Inflight service': 5, 'Cleanli
                'Departure Delay in Minutes': 1000, 'Arrival Delay in Minutes': 1000}

# create a DataFrame with the new data
custDf = pd.DataFrame(demoCustomer, index=[0])
# get the predicted satisfaction value for the new customer
prediction = tree.predict(custDf[predictors].values)
# convert the predicted value to a string
satisfaction = 'satisfied' if prediction[0] == 1 else 'neutral/dissatisfied'


print(f'The prediction for our demo customer is: {satisfaction}')
```

The prediction for our demo customer is: satisfied

Lastly, we have a confusion matrix to display our models accuracy.

```
In [668.. # Predictions here were from our final model before conclusion.
confusion = confusion_matrix(y_test, predictions)
# convert 0 to "neutral/dissatisfied", and 1 to "satisfied"
predictions = [0 if p==0 else 1 for p in predictions]
# convert 0 to "neutral/dissatisfied", and 1 to "satisfied"
y_test = [0 if y==0 else 1 for y in y_test]
# 1 if correct, 0 if incorrect
correct_predictions = [1 if p==t else 0 for p, t in zip(predictions, y_test)]
accuracy = sum(correct_predictions) / len(correct_predictions)
# Setup the matrix
fig, ax = plt.subplots(figsize=(6, 6))
ax.matshow(confusion, cmap=plt.cm.Blues, alpha=0.3)
for i in range(confusion.shape[0]):
    for j in range(confusion.shape[1]):
        ax.text(x=j, y=i, s=confusion[i, j], va='center', ha='center')
# Actually set the labels
tick_labels = ['Satisfied', 'Neutral/Dissatisfied']
ax.set_xticks([1, 0])
ax.set_yticks([1, 0])
ax.set_xticklabels(tick_labels)
ax.set_yticklabels(tick_labels)
# Axis & Display
plt.xlabel('Predicted Values')
```

```
plt.ylabel('Actual Values')  
plt.show()
```



Our matrix results show our 88% accuracy rating, as our previous model results had outputted. With 1541 false positives (predicted satisfied but actually dissatisfied) and 1481 false negatives (actual satisfied but predicted dissatisfied) for a total of 3022 incorrect predictions out of 25976 customers.

In conclusion, our group had learned a great deal from this class and project, we experimented with **kNN** and **Decision Trees** and found that due to the speed of Trees, we got a decent accurate result at a fraction of the time that kNN would calculate, even with less predictors. Specific to our project, while many of us can assume features like **Online boarding**, **Gate Location**, **Inflight Wifi**, **Type of Travel_Business travel**, or **Baggage Handling** would be quality predictors for ones enjoyment of a flight; those were also different than our initial hypothesis (Seat comfort, in-flight entertainment, cleanliness, and food & drinks). Even though we had found 10 features that were really accurate together, we were able to derive the best 5 features from our those and our overall list of 26 and determined which ones were stronger predictors and more relevant.

[Back to Top](#)